

Dog Breed Classifier

REVIEW

HISTORY

Meets Specifications

Congratulations! You've passed the project!

1. You did a great job by trying and testing out different things. We should always remember, we learn only when we experiment a lot and make mistakes :)
2. Dog breed classifier report by a Stanford student:
http://cs231n.stanford.edu/reports/2015/pdfs/fcdh_FinalReport.pdf
3. As you've already got familiarity with image classification, I'll recommend reading about object detection as well <https://towardsdatascience.com/r-cnn-fast-r-cnn-faster-r-cnn-yolo-object-detection-algorithms-36d53571365e>
4. Fun fact: In the real world we rarely train models from scratch. The already pre-trained models provide a good initialization for fine-tuning according to our tasks. Pre-trained weights are even used as weight initializers for training the whole network!

Good luck with your future projects!

Files Submitted

The submission includes all required, complete notebook files.

Step 1: Detect Humans

The submission returns the percentage of the first 100 images in the dog and human face datasets that

include a detected, human face.

Step 2: Detect Dogs

Use a pre-trained VGG16 Net to find the predicted class for a given image. Use this to complete a `dog_detector` function below that returns True if a dog is detected in an image (and False if not).

The submission returns the percentage of the first 100 images in the dog and human face datasets that include a detected dog.

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Write three separate data loaders for the training, validation, and test datasets of dog images. These images should be pre-processed to be of the correct size.

1. Some transforms that can be used for training images are
`transforms.RandomResizedCrop()`,
`transforms.RandomRotation()`,
`transforms.RandomHorizontalFlip()`
2. Shuffling required for training data, no shuffling required for validation and test. An article on randomness you may want to read <https://machinelearningmastery.com/randomness-in-machine-learning/>

Answer describes how the images were pre-processed and/or augmented.

1. You've correctly described the pre-processing step involving resizing, centre-cropping and augmentation.
2. Augmenting the dataset prevents overfitting as it introduces more variety to dataset in addition to increasing the dataset size. More on overfitting here <https://towardsdatascience.com/preventing-deep-neural-network-from-overfitting-953458db800a>

The submission specifies a CNN architecture.

1. Using at least 3 convolutional layers helps the model to be complex enough to accurately recognize dog breeds. Well done!
2. Pooling is a great way to reduce network parameters. It also introduces invariance in the model. You can read about uses of pooling and different pooling layers here <https://machinelearningmastery.com/pooling-layers-for-convolutional-neural-networks/>

3. Tip: Over the years, the research community has also tried an all convolutional network. Suggested paper <https://arxiv.org/abs/1412.6806>
4. You've used dropout which is an incredible technique to prevent overfitting of networks. Tip: Use dropout = 0.5 as that's what works well for everyone unless you have an experimental analysis or an intuitive reason for taking some other value.
5. Batch normalization can be used in CNNs to help deal with internal covariate shift. A useful article <https://towardsdatascience.com/batch-normalization-in-neural-networks-1ac91516821c>
6. Using ReLu is a good idea. The post here details some other activation functions used in neural networks <http://cs231n.github.io/neural-networks-1/>

Answer describes the reasoning behind the selection of layer types.

Choose appropriate loss and optimization functions for this classification task. Train the model for a number of epochs and save the "best" result.

1. You've used Adam which is one of the optimizers used by the deep learning community. You should definitely read about other optimizers here: <https://arxiv.org/pdf/1609.04747.pdf>
2. Another good post on optimizers: <https://ruder.io/optimizing-gradient-descent/index.html#adam>

The trained model attains at least 10% accuracy on the test set.

Well done!

Step 4: Create a CNN Using Transfer Learning

The submission specifies a model architecture that uses part of a pre-trained model.

1. You used VGG net. Good work! Other possible networks are ResNet variations, InceptionNet variations. Do try them as well if you have some time
2. Tip: Adding one or two fully connected layers of your own after the final layer of VGG and before the classification layer can improve the results you get.

The submission details why the chosen architecture is suitable for this classification task.

Train your model for a number of epochs and save the result with the lowest validation loss.

Accuracy on the test set is 60% or greater.

Well done!

The submission includes a function that takes a file path to an image as input and returns the dog breed that is predicted by the CNN.

Step 5: Write Your Algorithm

The submission uses the CNN from the previous step to detect dog breed. The submission has different output for each detected image type (dog, human, other) and provides either predicted actual (or resembling) dog breed.

Step 6: Test Your Algorithm

The submission tests at least 6 images, including at least two human and two dog images.

Submission provides at least three possible points of improvement for the classification algorithm.

1. Well analyzed! The listed points on improvement are indeed practical.
2. Tip: You may want to read on 'Ensembles of models'. More on it here <https://machinelearningmastery.com/ensemble-methods-for-deep-learning-neural-networks/>
3. Also, more augmentation techniques, larger ConvNets along with regularization techniques give better performing models.

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)

[Rate this project](#)

